

AKSHAR GOTHİ

San Francisco, CA | +1 (415) 747-2221 | akshargothi70@gmail.com | [linkedin.com/in/akshar-gothi/](https://www.linkedin.com/in/akshar-gothi/) | github.com/akshar27/

PROFESSIONAL SUMMARY

AI Engineer with 3+ years of experience building production LLM systems, AI agents, and scalable backend platforms. Skilled in LangGraph, OpenAI APIs, FastAPI, RAG, vector search, and cloud deployment. Proven track record shipping end-to-end AI products with memory, tool integrations, evaluation pipelines, and observability.

TECHNICAL SKILLS

- **Programming:** Python, JavaScript, TypeScript, Java, SQL
- **AI/ML:** LLMs, AI Agents, OpenAI APIs, LangGraph, LangChain, RAG, Prompt Engineering, Embeddings, Vector Search, Tool Calling, Transformers, PyTorch, Hugging Face, LangSmith, Fine-Tuning
- **Backend:** FastAPI, Node.js, Express, Django
- **Frontend:** React, Next.js
- **Databases:** PostgreSQL, MySQL, MongoDB, SQLite
- **Distributed Systems & Cloud:** AWS, GCP, Docker, CI/CD, Microservices, Railway, Vercel
- **ML Systems & Infra:** Model Deployment (FastAPI), Inference Optimization, Vector Databases (Qdrant/Pinecone), Model Evaluation, Experimentation

PROFESSIONAL EXPERIENCE

OpenAI (via Apex Systems) | AI Engineer (Contract) Jan 2026 - Present

- Partnered with ML and robotics engineers to evaluate real-world robotic AI behaviors, reducing failure rates by 20% through structured testing and root-cause analysis.
- Built repeatable testing workflows for embodied AI systems, accelerating debugging and model validation cycles.
- Executed data collection, task benchmarking, and failure analysis workflows to improve robotic AI system performance.
- Translated production system behaviors into actionable insights for engineering teams across hardware and software boundaries.

SFSU | Graduate Research Assistant Aug 2024 - Dec 2025

- Built HealthShare Insights - a production RAG platform analyzing 10K+ health social media posts using Qdrant, MiniLM, and Ollama.
- Researched emotion classification and sentiment analysis on COVID-era text data, building evaluation pipelines in Python.
- Implemented vector search pipelines with sentence-transformers and Qdrant, achieving sub-200ms retrieval latency.

Tatvasoft | Software Engineer Jul 2021 - Aug 2023

- Improved user activation and product discovery by 25% by optimizing search ranking systems and aligning backend logic with user behavior patterns.
- Conducted iterative experimentation on search ranking and API performance changes, analyzing engagement metrics to validate improvements in user activation and onboarding flows.
- Designed scalable microservice-based backend systems supporting millions of users with emphasis on fault tolerance, performance optimization, and high availability.
- Automated recurring operational workflows using Python and shell scripts, reducing manual workload by 30% and improving system efficiency.
- Worked closely with product teams to iterate on features based on user engagement data, ensuring alignment between technical solutions and business goals.

Reviewdale | Backend Engineer Intern Jul 2020 - Aug 2021

- Developed distributed microservices handling large-scale analytics workloads, improving scalability and service isolation.
- Designed and implemented secure RESTful APIs and integrated third-party services for payments and ad systems, ensuring reliability and data security.
- Optimized CI/CD pipelines using Docker and GCP, reducing deployment downtime and improving release velocity.

PROJECTS

AI Personal Assistant | <https://ai-personal-assistant-pi.vercel.app/> Mar 2026 - Apr 2026

- Built a production AI agent using LangGraph with multi-step tool-calling workflows across Gmail, Calendar, memory, and approval systems.
- Implemented context-aware email summarization, reply drafting, and approval-gated sensitive actions using human-in-the-loop workflows.
- Developed conflict-aware scheduling with persistent memory for user preferences such as concise emails and morning meetings.
- Implemented LLM evaluation pipelines measuring response quality, tool-calling accuracy, latency, and workflow reliability using LangSmith observability.
- Deployed full-stack application using FastAPI, Next.js, Railway, and Vercel with secure OAuth and production environment configs.

Shipify – Crowdsourced Package Delivery Platform | <https://shipfyy.com/> Jan 2025 - Jul 2025

- Designed matching algorithms and real-time filtering systems to improve activation success and reduce onboarding friction.
- Developed real-time messaging and tracking system using WebSockets, enabling seamless communication between users.

HealthShare Insights | <http://18.218.238.186/> Aug 2024 - Dec 2025

- Built a multimodal RAG chatbot over 10K+ health-related social media posts using Qdrant vector search, MiniLM embeddings, and Ollama llama3.2 for context-aware question answering.
- Designed end-to-end RAG pipeline: chunked PostgreSQL text data, generated embeddings via sentence-transformers, stored in Qdrant, and retrieved semantically relevant context for LLM responses.
- Built emotion analytics dashboards visualizing sentiment trends, word frequency, and source comparisons across COVID-era health data.
- Deployed full-stack application on AWS EC2 with Node.js backend, PostgreSQL, and local Qdrant + Ollama inference stack.

YouDescribeX | Accessibility AI Platform | <https://ydx.youdescribe.org/home> Jul 2024 - Dec 2024

- Built a multimodal RAG-based AI platform generating contextual video captions for visually impaired users using Qdrant vector search and LLM reasoning.
- Fine-tuned Llama 2 using LoRA/QLoRA on domain-specific video data, improving response quality by 30% and reducing hallucinations.
- Reduced inference latency by 50% using optimized pipelines, caching, and scalable inference workflows.
- Built evaluation pipelines measuring model accuracy, retrieval quality, latency, and response quality.

EDUCATION

San Francisco State University | *Master of Science, Computer Science* (GPA: 3.63/4) San Francisco, California, USA Aug 2023 - Dec 2025

Chandubhai S Patel Institute of Technology | *Bachelor of Technology, Information Technology* (GPA: 3.5/4) Gujarat, India Jul 2017 - May 2021